

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt

1 from matplotlib.colors import ListedColormap
2 cmap_bold = ListedColormap(['#FF0000', '#0000FF'])
3 cmap_light = ListedColormap(['#FFBBBB', '#BBBBFF'])

1 # Step 1: Load the dataset
2 data = pd.read_csv("/content/Churn_Modelling - Churn_Modelling.csv")

1 # Step 2: Data preprocessing
2 # Encode categorical variables
3 data['Gender'] = data['Gender'].map({'Female': 0, 'Male': 1}) # Encode Gender
4 data = pd.get_dummies(data, columns=['Geography']) # One-hot encoding for Geography
5
6 # Select features and target variable
7 X = data[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
8           'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Gender',
9           'Geography_France', 'Geography_Germany', 'Geography_Spain']].values
10 y = data['Exited'].values.reshape(-1, 1) # Reshape y for matrix operations
11
12 # Normalize features
13 X = X.astype(np.float64)
14 X = (X - np.mean(X, axis=0)) / np.std(X, axis=0)
15


1 def sigmoid(h):
2     return 1/(1+np.exp(-h))
3
4 def bin_cross_entropy(y,p_hat):
5     return -(1/len(y))*np.sum(y*np.log(p_hat)+(1-y)*np.log(1-p_hat))
6
7 def accuracy(y,y_hat):
8     return np.mean(y==y_hat)

1 class LogisticRegression():
2     def __init__(self,thresh=0.5):
3         self.thresh=thresh
4         self.W = None
5         self.b = None
6
7     def fit(self,X,y,eta=1e-3,epochs=1e3,show_curve=True):
8         epochs = int(epochs)
9         N,D = X.shape
10
11         #Initailize Weights and Biases
12         self.W = np.random.randn(D)
13         self.b = np.random.randn(1)
14         J = np.zeros(epochs)
15
16         #Stochastic Gradient Descent
17         for epoch in range(epochs):
18             p_hat = self.__forward__(X)
19             J[epoch] = bin_cross_entropy(y,p_hat)
20
21             #WeightUpdate Rules
22             error = (p_hat - y.ravel())
23             self.W -= eta*(1/N)*X.T@error
24             self.b -= eta*(1/N)*np.sum(error)
25         if show_curve:
26             plt.figure()
27             plt.plot(J)
28             plt.xlabel("epochs")
29             plt.ylabel("$\mathcal{J}$")
30             plt.title("Training Curve")
31             plt.show()
32
33     def __forward__(self,X):
34         return sigmoid(X@self.W+self.b)
35
36     def predict(self,X):
37         return (self.__forward__(X)>= self.thresh).astype(np.int32)

```

```
return (self._forward(X) - self._loss(y_hat, y)).astype(np.float32)
```

```
1 log_reg = LogisticRegression()
2 log_reg.fit(X,y, epochs =1e4, eta=1e-2, show_curve=True)
3 y_hat = log_reg.predict(X)
4 print(f"Training Accuracy: {accuracy(y,y_hat):0.4f}")
5 print(log_reg.W,log_reg.b)
```



Training Accuracy: 0.7528

```
[-0.06457684  0.76219181 -0.04609227  0.16450587 -0.05903227 -0.02034948
 -0.53715938  0.02763329 -0.26304535 -0.35818206  0.02522075 -0.29403544] [-1.65302541
```

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.